MapReduce Job for Data Processing on Google Cloud Platform

SHARMELE SOMU

Contents

CREATION OF CLUSTERS

STEP 1:

To create a dataproc cluster, open Google cloud console, type ‘dataproc’ on the search bar and select Dataproc

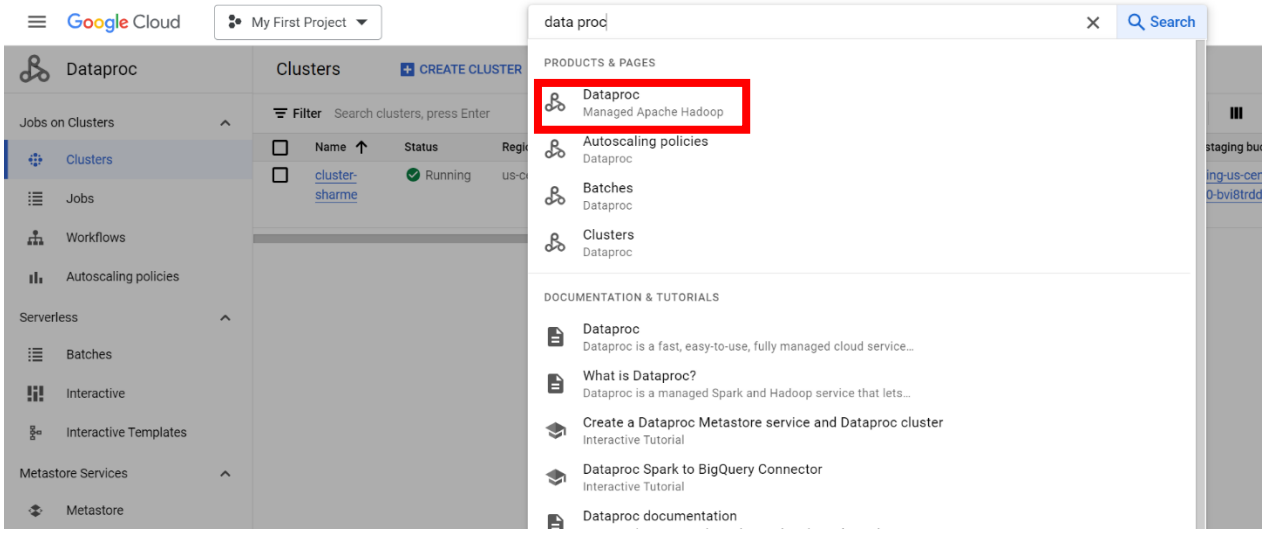

STEP 2:

Click on create cluster to create the cluster.

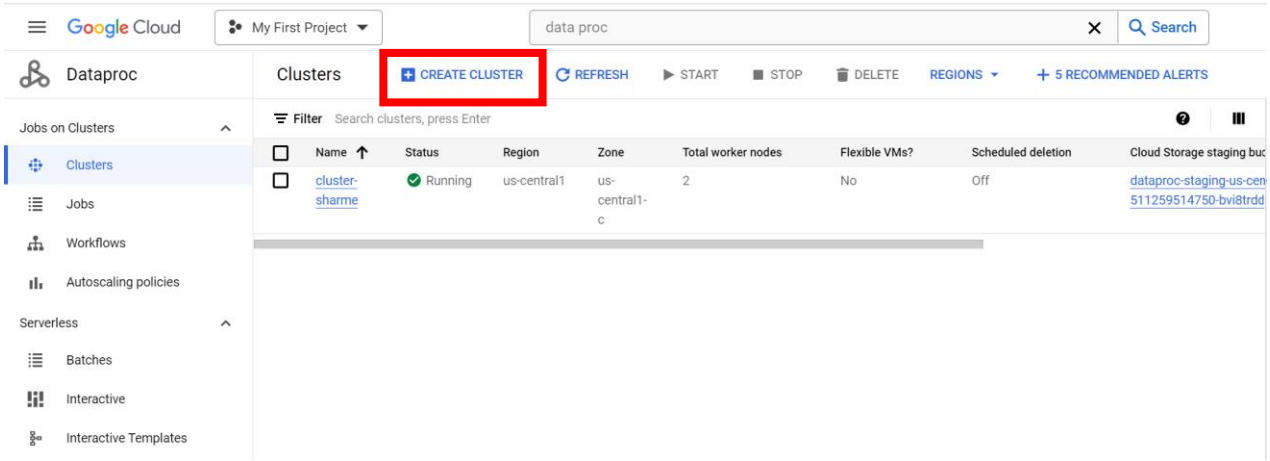

STEP 3:

Click on create to create the cluster on compute engine.

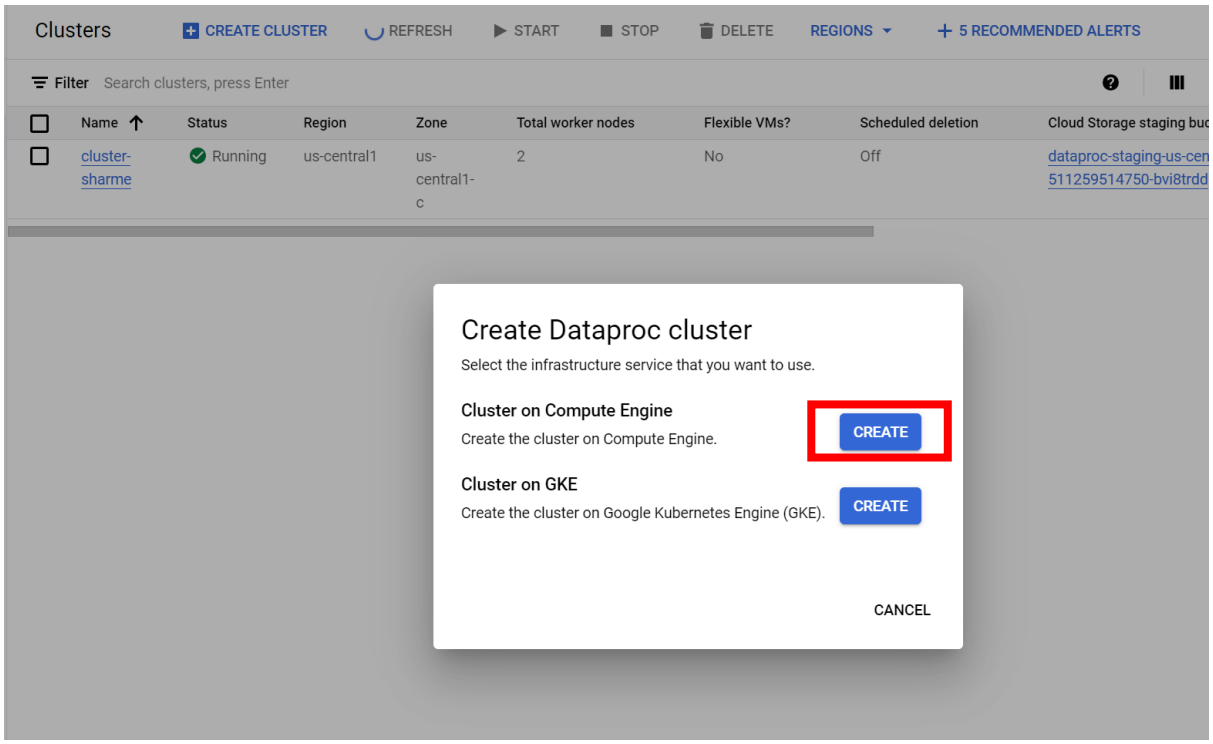

STEP 4:

Next step is to set up the cluster. Mention the name of the cluster.

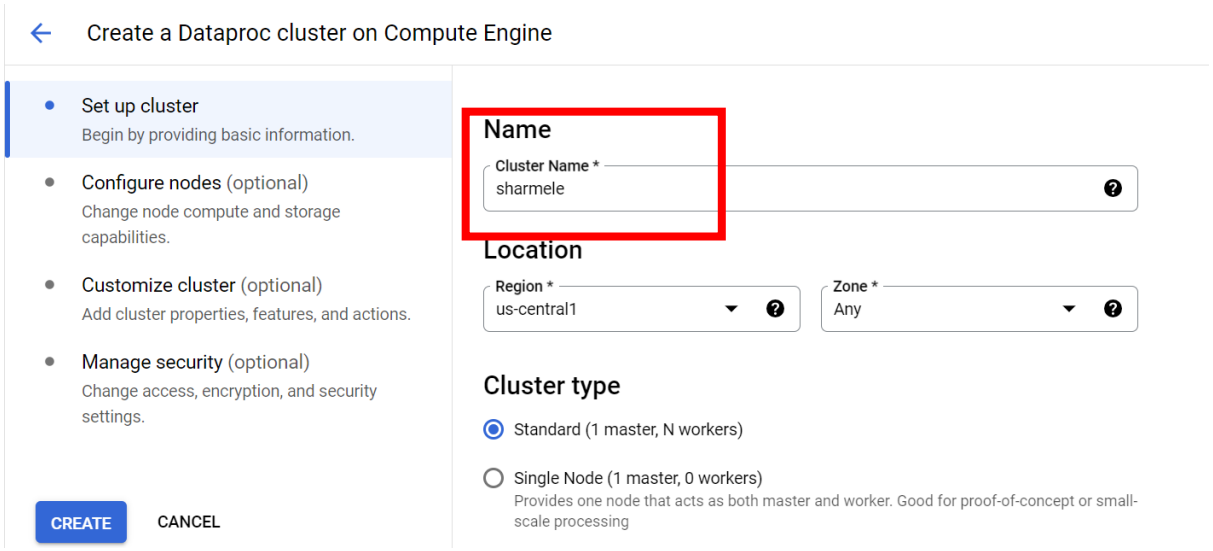

STEP 5:

Select the name of operating system and version of operating system, Hadoop and spark.

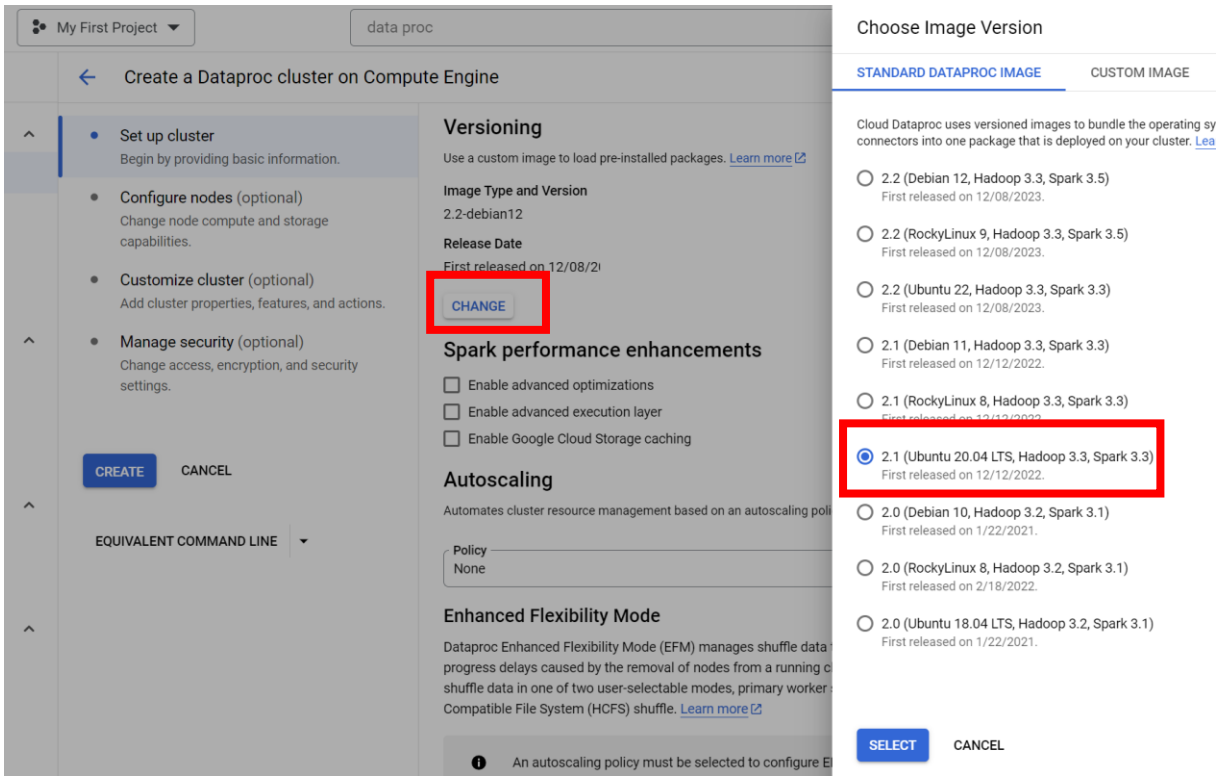

STEP 6:

In configure nodes, setup the manager and worker nodes. Select the series type, machine type and disk size for the manager node.

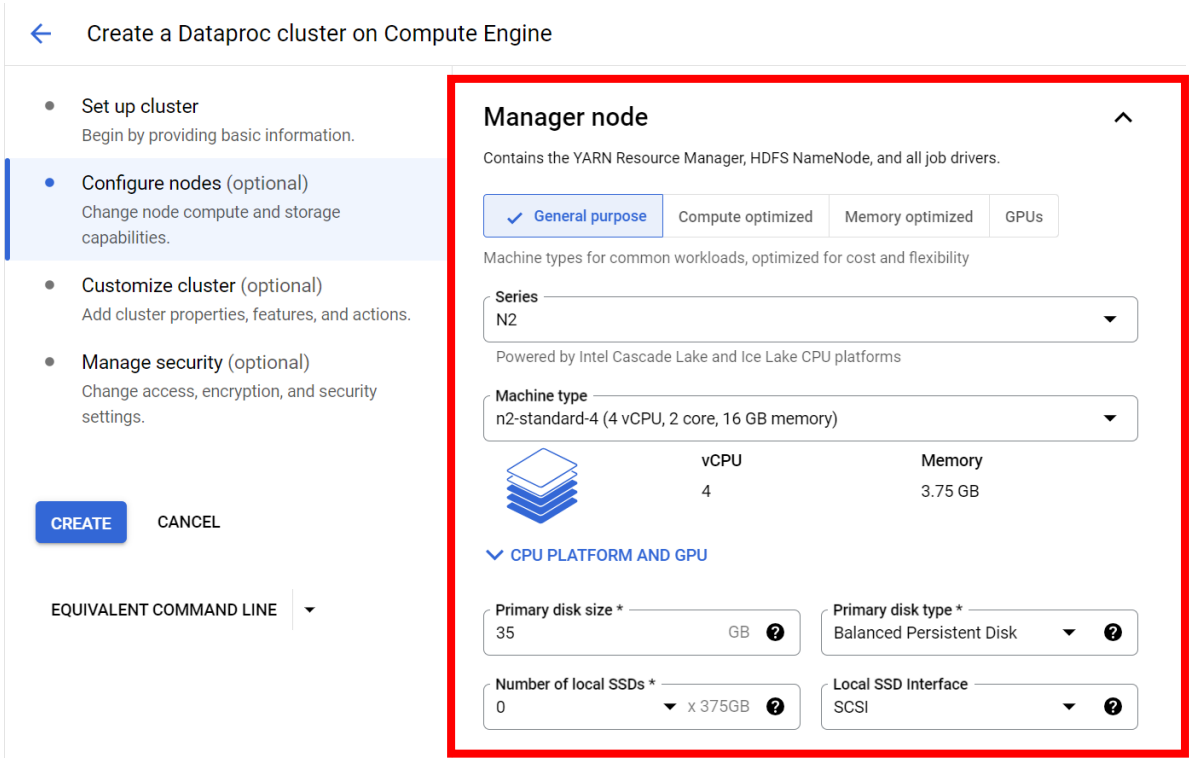

STEP 7:

Select the series type, machine type and disk size for the worker node and click on create to create the cluster with 1 manager/name node and 2 worker nodes.

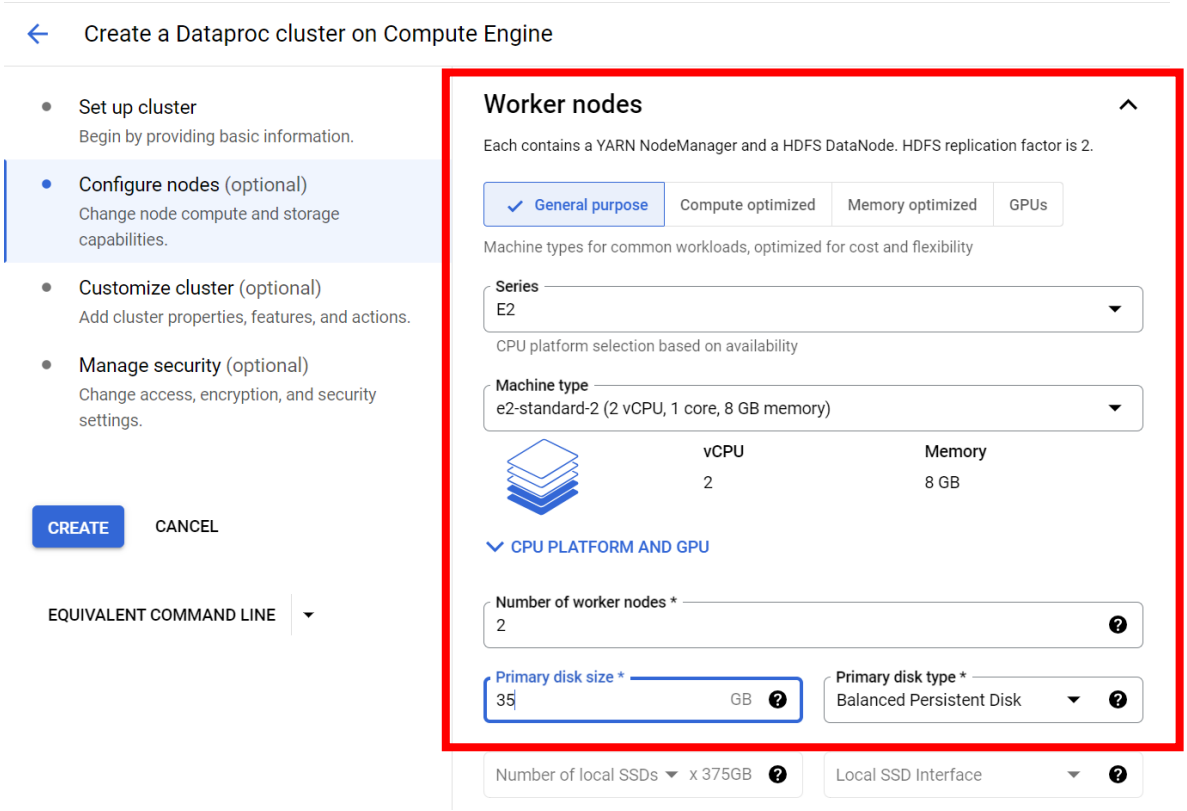

STEP 8:

Click on the cluster to view the VM instances.

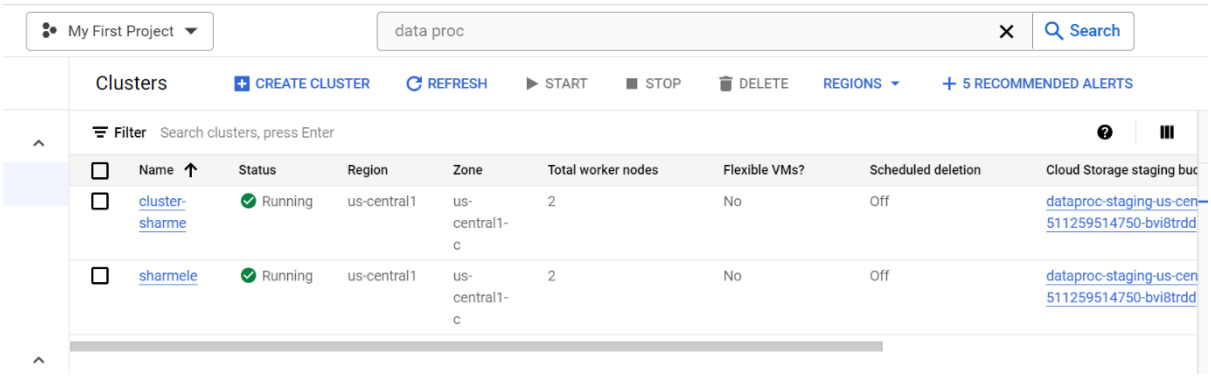

STEP 9:

Click on SSH to connect to the manager node. It opens the terminal window directly on the browser allowing users to interact with the VM's command line interface.

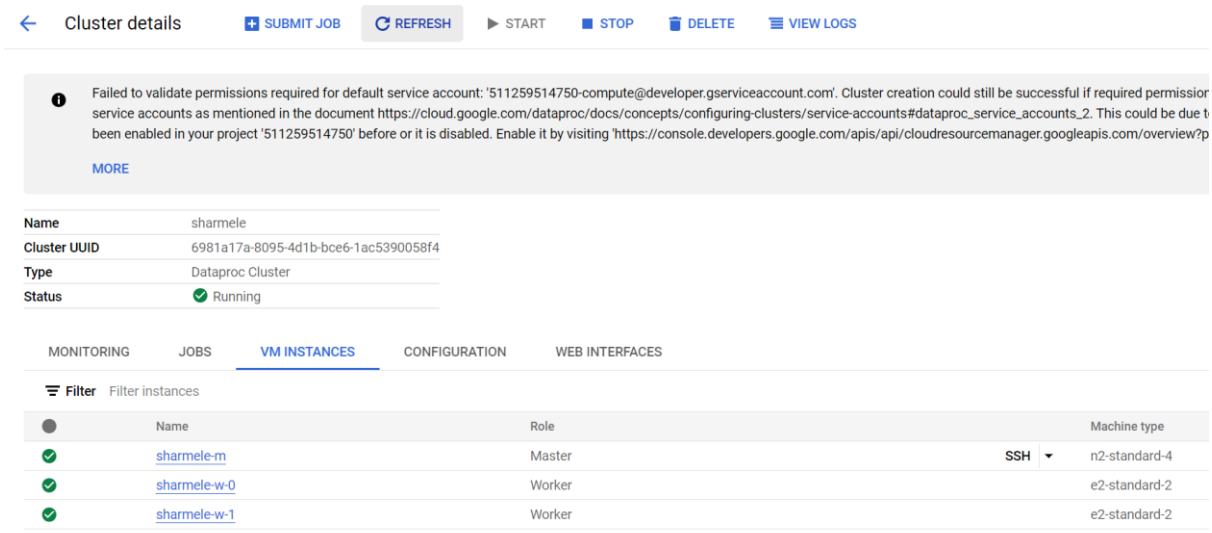

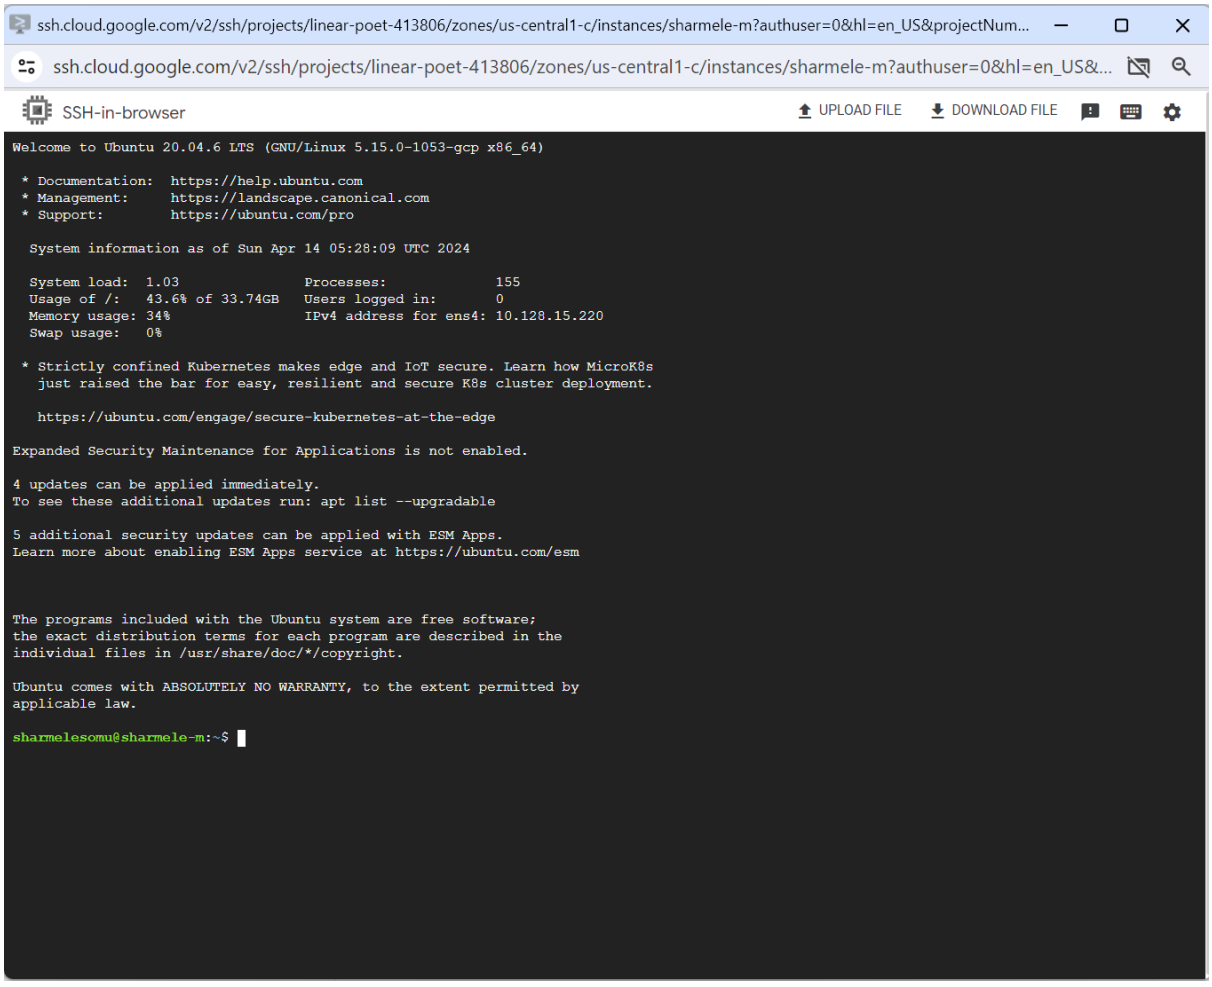

TASK A: EXECUTING THE MAPPER AND REDUCER CODE ON THE NAME NODE IN GCP

The objective is to join two tables TA.csv and TB.csv and display the student records for those students whose date of birth is greater than '01/01/1997'.

MAPPER EXPLANATION

The mapper code reads the input from the standard input, extracts fields by using ' , ' as the delimiter, checks if the fields are from table A or table B. If the fields are from table A, the records are assigned value 1. If the fields are from table B the records are assigned value 2. Then the output of the mapper is displayed as a key-value pair. Joining of tables is done at the reducer side.

REFERENCES

I have used the following references and ChatGPT to generate the code for the mapper and reducer.

https://www.youtube.com/watch?v=ai0E4ovoA5k&t=151s
https://www.edureka.co/blog/mapreduce-example-reduce-side-join/
https://ars.els-cdn.com/content/image/1-s2.0-S1319157820303694-gr5_lrg.jpg

MAPPER CODE – mapper.py

```
#!/usr/bin/env python
import sys
value=0
# reads input from standard input
for line in sys.stdin:
# extracts fields by using ',' as a delimiter
  fields = line.strip().split(',')
# Assigns fields to student_id, name and dob
  student_id = fields[0]
  name = fields[1]
  dob = fields[2]
# In both tables, when StudentId is encounter, the header is skipped
  if fields[0] == 'StudentId':
    continue
# Assigning value 1 for records from table A to generate a composite key
  elif fields[1] in ['Alice','Tom','John']:
    value = 1
# Assigning value 2 for records from table B to generate a composite key
  else:
    value = 2
# print the mapper output - key and values
  print(student_id,value,[name,dob])
```

REDUCER EXPLANATION

The reducer code reads the data from standard input, splits the data with the table number. If the table number is ‘1’, the code converts the date from string to date time format and stores the data in student_info_tableA. If the table number is ‘2’, the code appends the course details to student_id. Then it checks if student_id from student_info_tableA is present in student_info_tableB and joines the tables. There is a counter logic which counts the number of occurrences of each student_id. Then the student records are filtered to check if date of birth is than ‘01/01/1997’ and the resulting records are printed.

REDUCER CODE – reducer1.py

```
import sys
from datetime import datetime
# Dictionary to store information for each student ID
student_info_tableA = {}
student_info_tableB = {}
prev_student_id = None
counter = 0
x = 0
student_records = {}
# Read input lines from standard input
for line in sys.stdin:
  # Split the line into its components: student ID, table number, and the rest of the data
  student_id, table_number, rest = line.strip().split(' ', 2)
  # If the table number is '1', update student_info_tableA with the student's name and DOB
  if table_number == '1':
    name, dob_str = eval(rest)
    dob = datetime.strptime(dob_str, '%m/%d/%Y').date()
    student_info_tableA.setdefault(student_id, {'name': name, 'dob': dob})
  # If the table number is '2', update student_info_tableB with the student's course data
  elif table_number == '2':
    course_id, grade = eval(rest)
    student_info_tableB.setdefault(student_id, []).append((course_id, grade))
# Perform the join operation and filter records
for student_id in student_info_tableA:
  # Check if the student is present in both tables
  if student_id in student_info_tableB:
    # Combine the information from both tables
    for course_data in student_info_tableB[student_id]:
      name = student_info_tableA[student_id]['name']
      dob = student_info_tableA[student_id]['dob']
        # counter value to tag each student record
      if prev_student_id == student_id:
        counter = counter
        x=counter
      else:
        counter += 1
```

```
            prev_student_id = student_id
        # filters those student records for which date is above '01/01/1995'
         if dob > datetime(1995, 1, 1).date():  # Filter based on date of birth
            student_records.setdefault(student_id, []).append([name, dob_str, course_data[0],
course_data[1]])
      for student_id, records in student_records.items():
         counter = 0  # Initialize counter for each student ID
         for record in records:
            counter += 1  # Increment the counter for each record
            print(f'{student_id}\t{counter}\t{record}') # print the student records after filtering dob
```

EXECUTING THE MAPPER AND REDUCER ON THE NAME NODE IN GCP

STEP 1: Upload the csv files – TA.csv, TB.csv, python files - mapper and reducer on the name node.

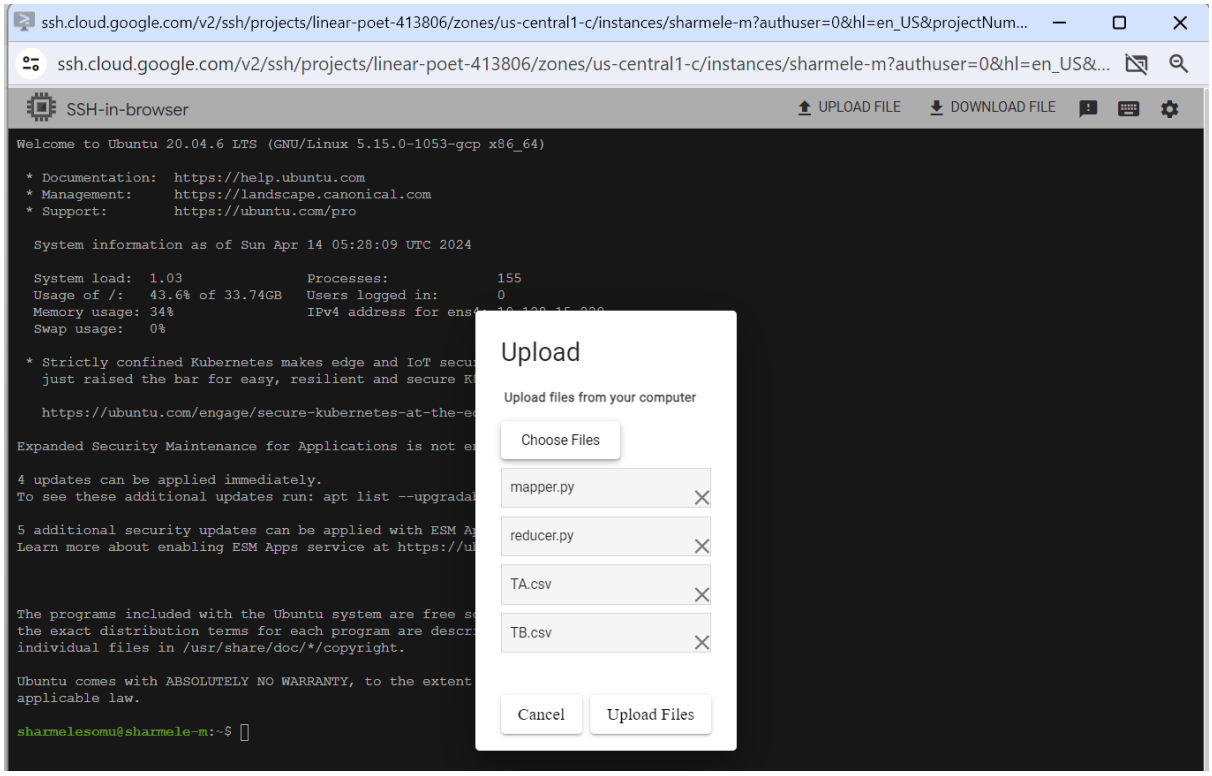

STEP 2: use the command 'ls' to check the list of loaded files.

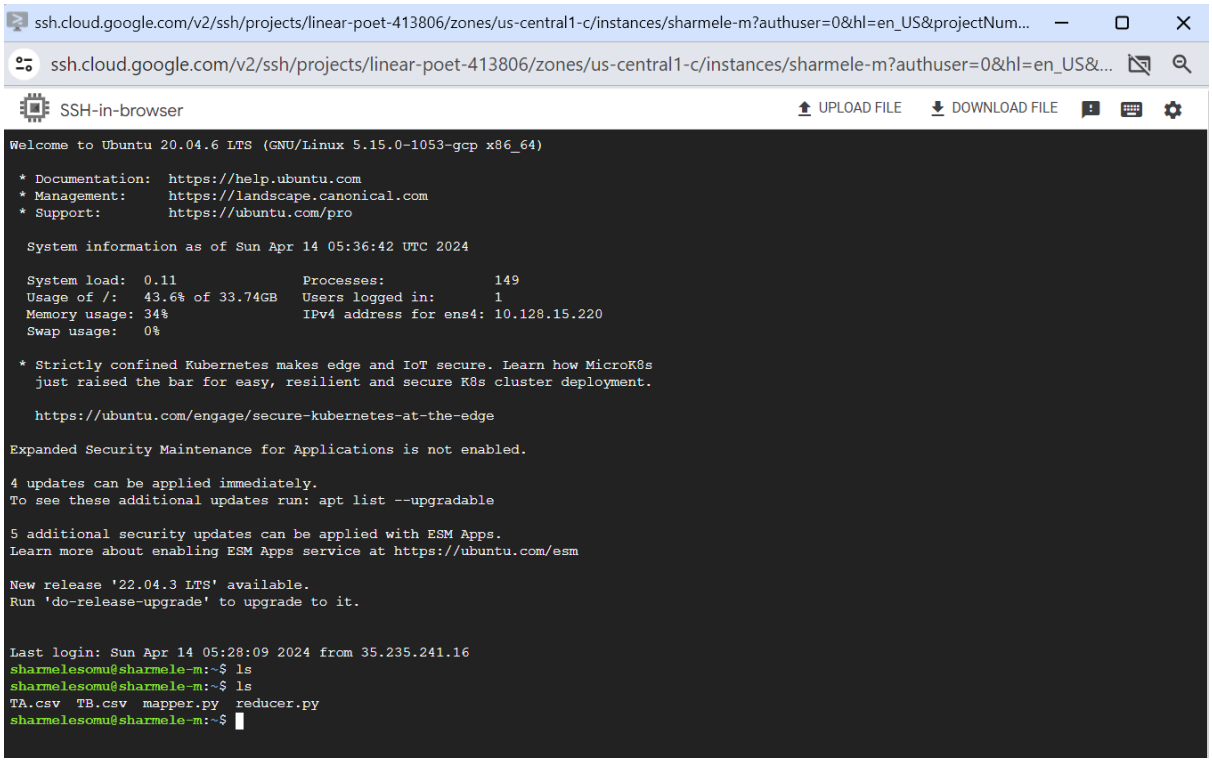

STEP 3: Display the content of Table 1 and Table 2 using **'cat'** command.

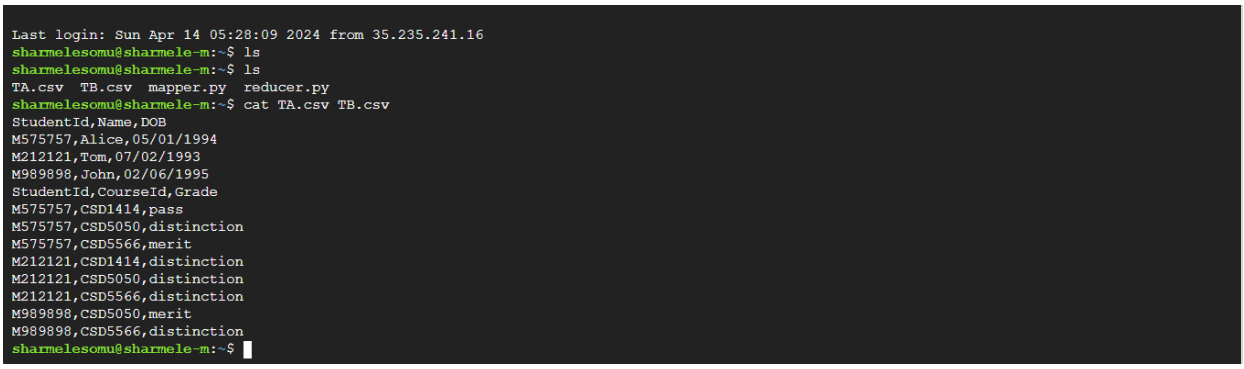

STEP 4: Execute the mapper code – mapper.py by using the command python mapper.py. The output from above cat command is the input to the mapper. We use the '|' symbol to input the output from one command to another command

cat TA.csv TB.csv | python mapper.py

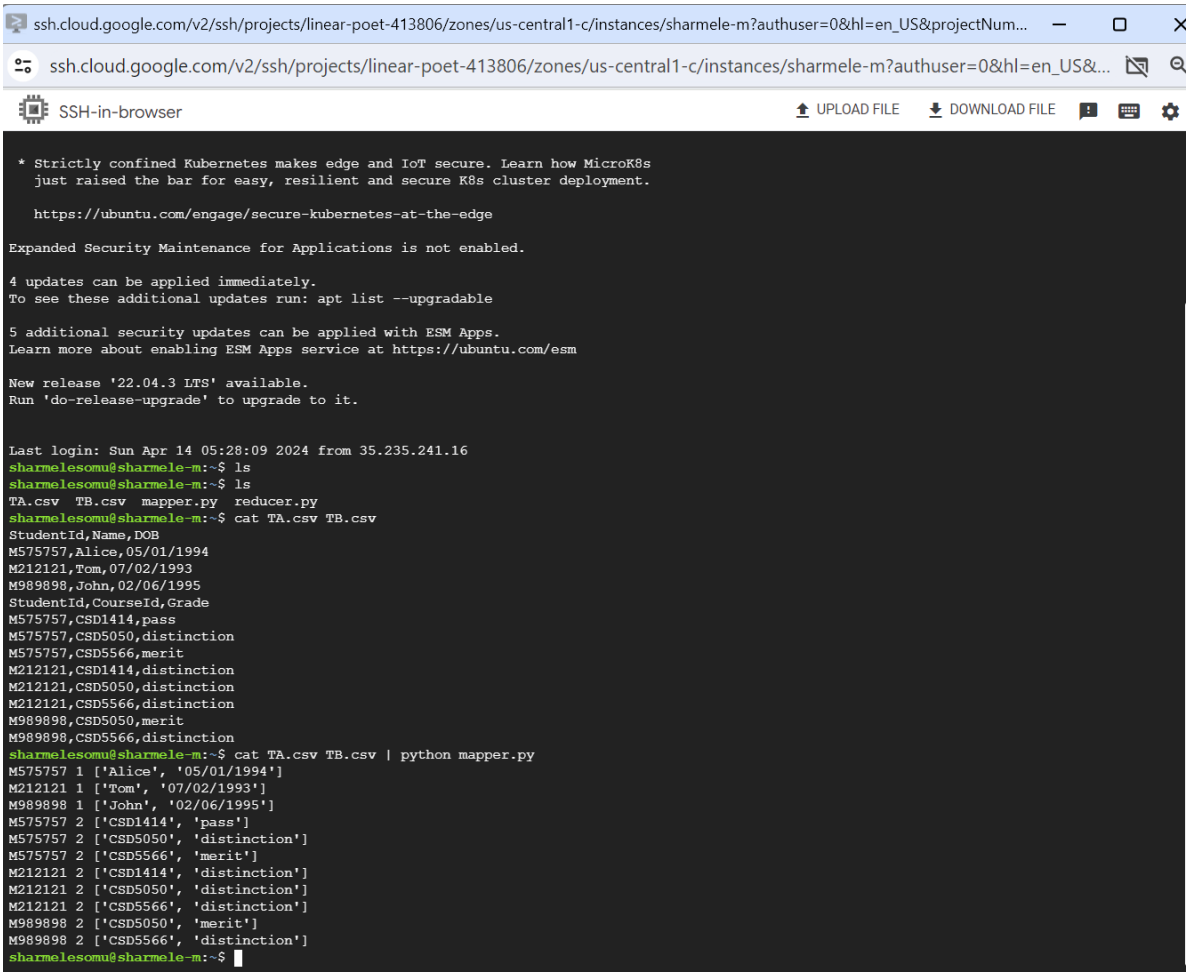

STEP 5: Sort the output from the mapper.

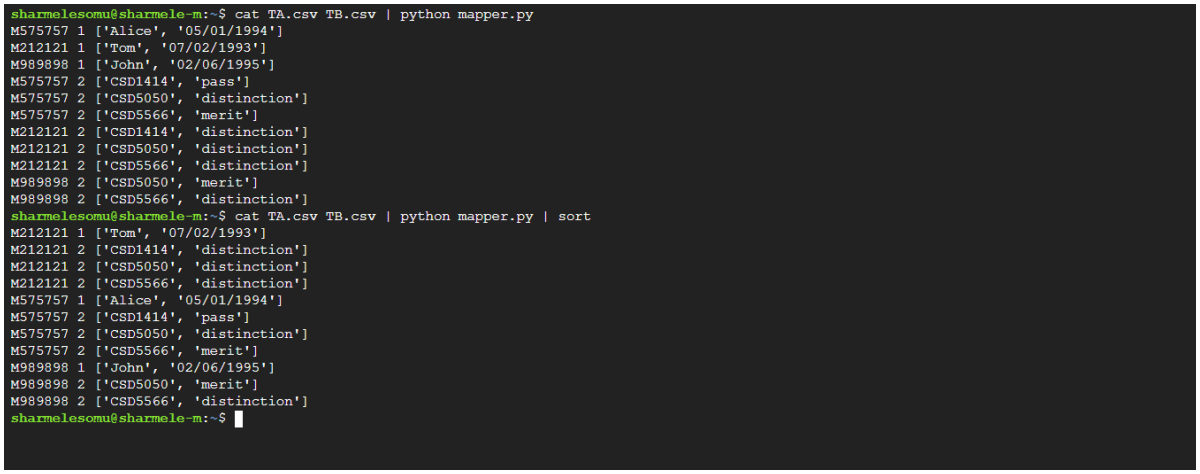

STEP 5: Executing the reducer code – reducer1.py by using the command python reducer.py. The output from the mapper is given as an input to the reducer.

cat TA.csv TB.csv | python mapper.py |sort | python reducer1.py

```
sharmelesomu@sharmele-m:~$ cat TA.csv TB.csv | python mapper.py
M575757 1 ['Alice', '05/01/1994']
M212121 1 ['Tom', '07/02/1993']
M989898 1 ['John', '02/06/1995']
M575757 2 ['CSD1414', 'pass']
M575757 2 ['CSD5050', 'distinction']
M575757 2 ['CSD5566', 'merit']
M212121 2 ['CSD1414', 'distinction']
M212121 2 ['CSD5050', 'distinction']
M212121 2 ['CSD5566', 'distinction']
M989898 2 ['CSD5050', 'merit']
M989898 2 ['CSD5566', 'distinction']
sharmelesomu@sharmele-m:~$ cat TA.csv TB.csv | python mapper.py| sort
M212121 1 ['Tom', '07/02/1993']
M212121 2 ['CSD1414', 'distinction']
M212121 2 ['CSD5050', 'distinction']
M212121 2 ['CSD5566', 'distinction']
M575757 1 ['Alice', '05/01/1994']
M575757 2 ['CSD1414', 'pass']
M575757 2 ['CSD5050', 'distinction']
M575757 2 ['CSD5566', 'merit']
M989898 1 ['John', '02/06/1995']
M989898 2 ['CSD5050', 'merit']
M989898 2 ['CSD5566', 'distinction']
sharmelesomu@sharmele-m:~$ cat TA.csv TB.csv | python mapper.py| sort| python reducer1.py
M989898 1       ['John', '02/06/1995', 'CSD5050', 'merit']
M989898 2       ['John', '02/06/1995', 'CSD5566', 'distinction']
sharmelesomu@sharmele-m:~$
```

TASK B: RESULTS OF THE MAPPER AND REDUCER ON NAME NODE

RESULT OF MAPPER

Mapper combines both tables and displays in key-value format

```
sharmelesomu@sharmele-m:~$ ls
TA.csv  TB.csv  mapper.py  red_trial_dict1.py  reducer.py  reducer1.py
sharmelesomu@sharmele-m:~$ cat TA.csv TB.csv
StudentId,Name,DOB
M575757,Alice,05/01/1994
M212121,Tom,07/02/1993
M989898,John,02/06/1995
StudentId,CourseId,Grade
M575757,CSD1414,pass
M575757,CSD5050,distinction
M575757,CSD5566,merit
M212121,CSD1414,distinction
M212121,CSD5050,distinction
M212121,CSD5566,distinction
M989898,CSD5050,merit
M989898,CSD5566,distinction
sharmelesomu@sharmele-m:~$ cat TA.csv TB.csv | python mapper.py
M575757 1 ['Alice', '05/01/1994']
M212121 1 ['Tom', '07/02/1993']
M989898 1 ['John', '02/06/1995']
M575757 2 ['CSD1414', 'pass']
M575757 2 ['CSD5050', 'distinction']
M575757 2 ['CSD5566', 'merit']
M212121 2 ['CSD1414', 'distinction']
M212121 2 ['CSD5050', 'distinction']
M212121 2 ['CSD5566', 'distinction']
M989898 2 ['CSD5050', 'merit']
M989898 2 ['CSD5566', 'distinction']
sharmelesomu@sharmele-m:~$
```

RESULT OF REDUCER

Reducer joins the tables and displays those records for which date of birth is greater than ’01/01/1995’

```
sharmelesomu@sharmele-m:~$ cat TA.csv TB.csv | python mapper.py| sort
M212121 1 ['Tom', '07/02/1993']
M212121 2 ['CSD1414', 'distinction']
M212121 2 ['CSD5050', 'distinction']
M212121 2 ['CSD5566', 'distinction']
M575757 1 ['Alice', '05/01/1994']
M575757 2 ['CSD1414', 'pass']
M575757 2 ['CSD5050', 'distinction']
M575757 2 ['CSD5566', 'merit']
M989898 1 ['John', '02/06/1995']
M989898 2 ['CSD5050', 'merit']
M989898 2 ['CSD5566', 'distinction']
sharmelesomu@sharmele-m:~$ cat TA.csv TB.csv | python mapper.py| sort| python reducer1.py
M989898 1       ['John', '02/06/1995', 'CSD5050', 'merit']
M989898 2       ['John', '02/06/1995', 'CSD5566', 'distinction']
sharmelesomu@sharmele-m:~$
```

TASK C: EXECUTING THE MAPPER AND REDUCER ON HADOOP STREAMING SERVICE

PREPARING THE HADOOP ENVIRONMENT

STEP 1: Create default folder in the Hadoop filesystem.

Command: Hadoop fs -mkdir /user/sharmelesomu

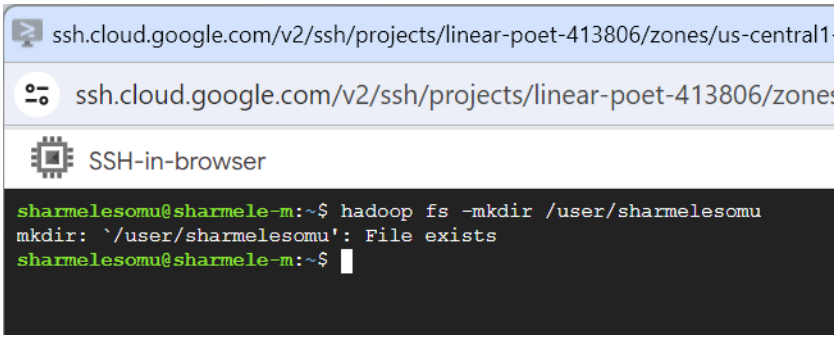

STEP 2: Store the input files on the Hadoop default folder

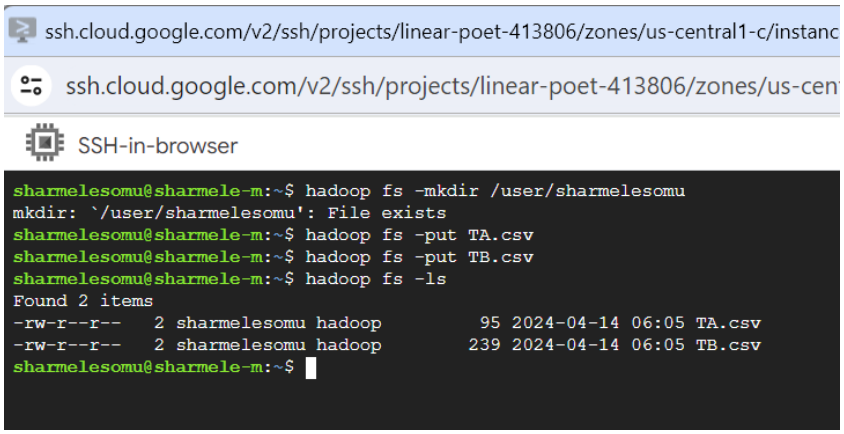

STEP 3: Execute the mapper using Hadoop streaming service.

Below command is used to execute the mapper on the Hadoop environment.

**hadoop jar /usr/lib/hadoop/hadoop-streaming-3.3.6.jar **

**-files /home/sharmelesomu/mapper.py -mapper "/usr/bin/python3 mapper.py" **

**-input /user/sharmelesomu/TA.csv,/user/sharmelesomu/TB.csv **

-output /user/sharmelesomu/output11

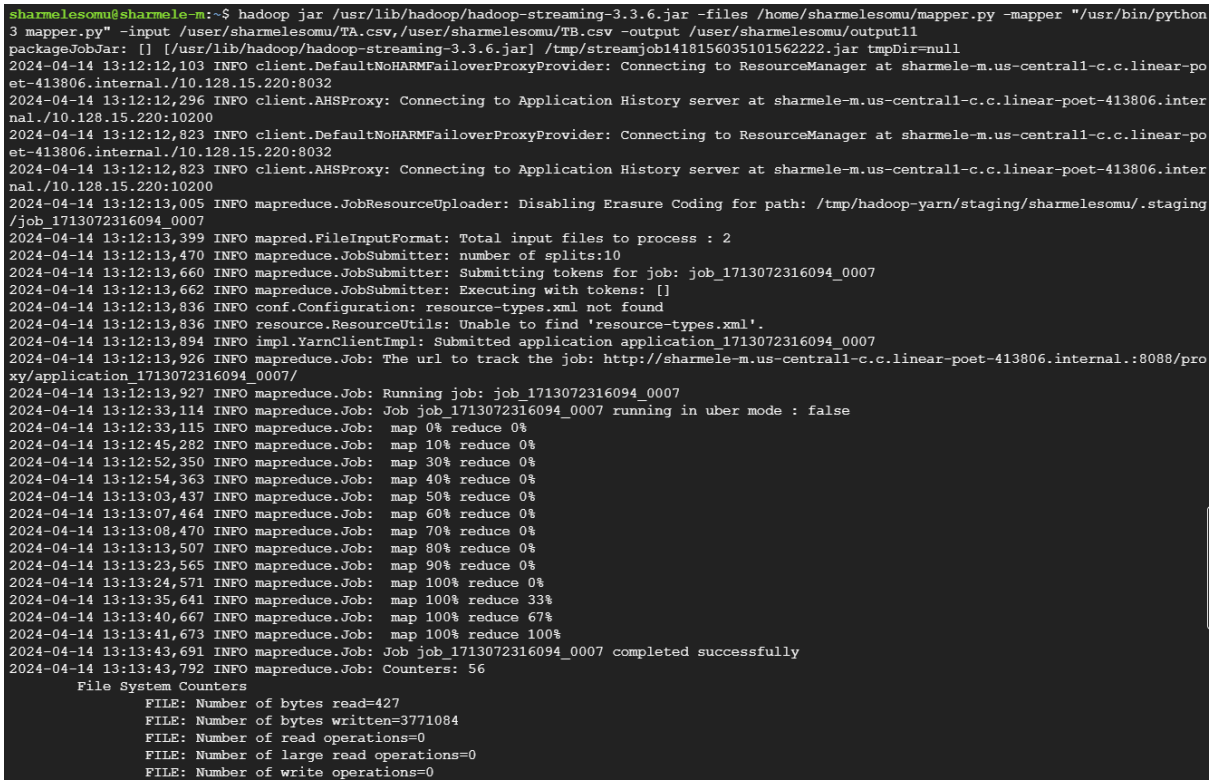

STEP 4: Execute the reducer using Hadoop streaming service.

Below command is used for executing the reducer in the Hadoop environment

hadoop jar /usr/lib/hadoop/hadoop-streaming-3.3.6.jar

-D mapreduce.job.reduces=1

-files /home/sharmelesomu/mapper.py,/home/sharmelesomu/reducer1.py

-mapper "/usr/bin/python3 mapper.py" -reducer "/usr/bin/python3 reducer1.py"

-input /user/sharmelesomu/TA.csv,/user/sharmelesomu/TB.csv

-output /user/sharmelesomu/output16

```
sharmelesomu@sharmele-m:~$ hadoop jar /usr/lib/hadoop/hadoop-streaming-3.3.6.jar -D mapreduce.job.reduces=1 -files /home/sharmelesomu/mapper.
py,/home/sharmelesomu/reducer1.py -mapper "/usr/bin/python3 mapper.py" -reducer "/usr/bin/python3 reducer1.py" -input /user/sharmelesomu/TA.c
sv,/user/sharmelesomu/TB.csv -output /user/sharmelesomu/output16
packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-3.3.6.jar] /tmp/streamjob13952913673159168455.jar tmpDir=null
2024-04-14 16:12:16,297 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at sharmele-m.us-central1-c.c.linear-po
et-413806.internal./10.128.15.222:8032
2024-04-14 16:12:16,497 INFO client.AHSProxy: Connecting to Application History server at sharmele-m.us-central1-c.c.linear-poet-413806.inter
nal./10.128.15.222:10200
2024-04-14 16:12:16,929 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at sharmele-m.us-central1-c.c.linear-po
et-413806.internal./10.128.15.222:8032
2024-04-14 16:12:16,929 INFO client.AHSProxy: Connecting to Application History server at sharmele-m.us-central1-c.c.linear-poet-413806.inter
nal./10.128.15.222:10200
2024-04-14 16:12:17,127 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/sharmelesomu/.staging
/job_1713104652195_0004
2024-04-14 16:12:17,538 INFO mapred.FileInputFormat: Total input files to process : 2
2024-04-14 16:12:17,638 INFO mapreduce.JobSubmitter: number of splits:10
2024-04-14 16:12:17,885 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1713104652195_0004
2024-04-14 16:12:17,885 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-04-14 16:12:18,075 INFO conf.Configuration: resource-types.xml not found
2024-04-14 16:12:18,076 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-04-14 16:12:18,154 INFO impl.YarnClientImpl: Submitted application application_1713104652195_0004
2024-04-14 16:12:18,202 INFO mapreduce.Job: The url to track the job: http://sharmele-m.us-central1-c.c.linear-poet-413806.internal.:8088/pro
xy/application_1713104652195_0004/
2024-04-14 16:12:18,203 INFO mapreduce.Job: Running job: job_1713104652195_0004
2024-04-14 16:12:33,398 INFO mapreduce.Job: Job job_1713104652195_0004 running in uber mode : false
2024-04-14 16:12:33,399 INFO mapreduce.Job:  map 0% reduce 0%
2024-04-14 16:12:44,516 INFO mapreduce.Job:  map 10% reduce 0%
2024-04-14 16:12:51,565 INFO mapreduce.Job:  map 30% reduce 0%
2024-04-14 16:12:53,578 INFO mapreduce.Job:  map 40% reduce 0%
2024-04-14 16:13:02,633 INFO mapreduce.Job:  map 50% reduce 0%
2024-04-14 16:13:06,660 INFO mapreduce.Job:  map 70% reduce 0%
2024-04-14 16:13:11,693 INFO mapreduce.Job:  map 80% reduce 0%
2024-04-14 16:13:21,746 INFO mapreduce.Job:  map 90% reduce 0%
2024-04-14 16:13:22,751 INFO mapreduce.Job:  map 100% reduce 0%
2024-04-14 16:13:33,833 INFO mapreduce.Job:  map 100% reduce 100%
2024-04-14 16:13:35,851 INFO mapreduce.Job: Job job_1713104652195_0004 completed successfully
2024-04-14 16:13:35,965 INFO mapreduce.Job: Counters: 55
        File System Counters
                FILE: Number of bytes read=415
                FILE: Number of bytes written=3201598
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
```

TASK D: RESULT OF THE MAPPER AND REDUCER USING HADOOP STREAMING SERVICE

RESULT OF THE MAPPER IN HADOOP ENVIRONMENT

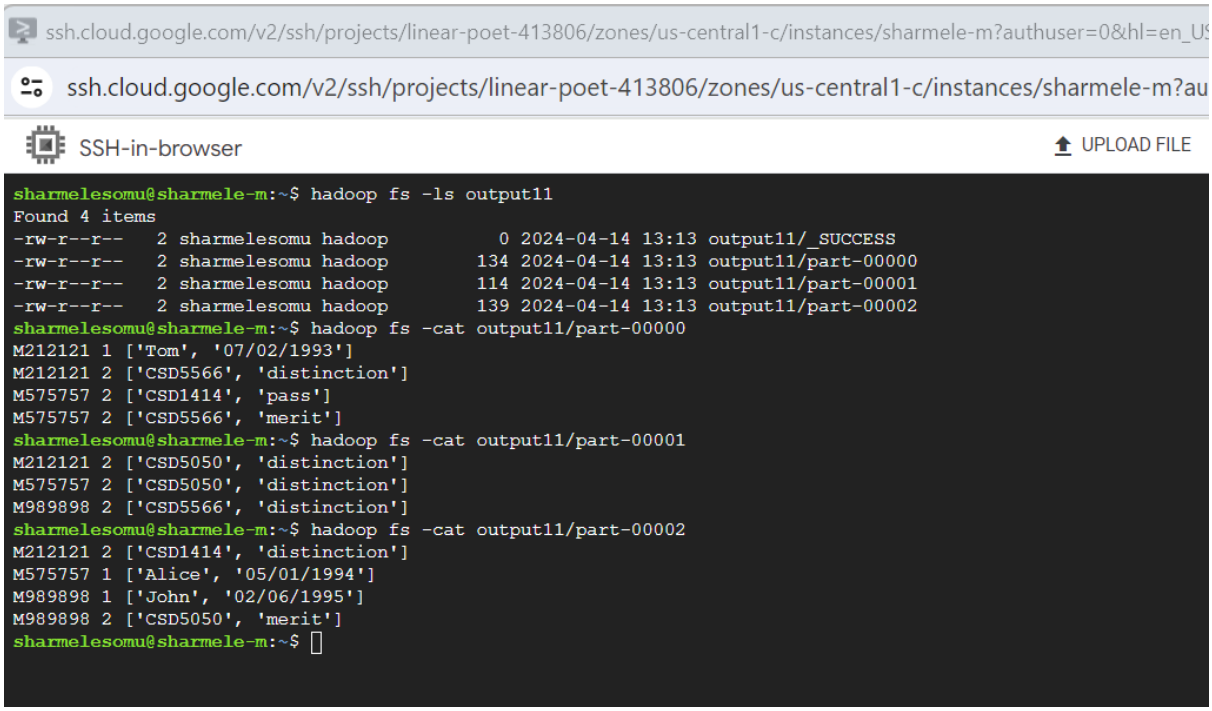

RESULT OF THE REDUCER IN HADOOP ENVIRONMENT

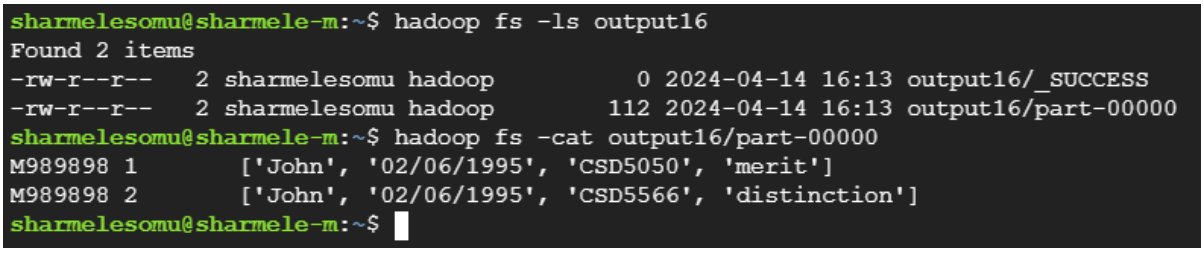

To view the number of mappers and reducers in the Google cloud platform. Navigate to the created cluster, select web interfaces, and click on Mapreduce Job History.

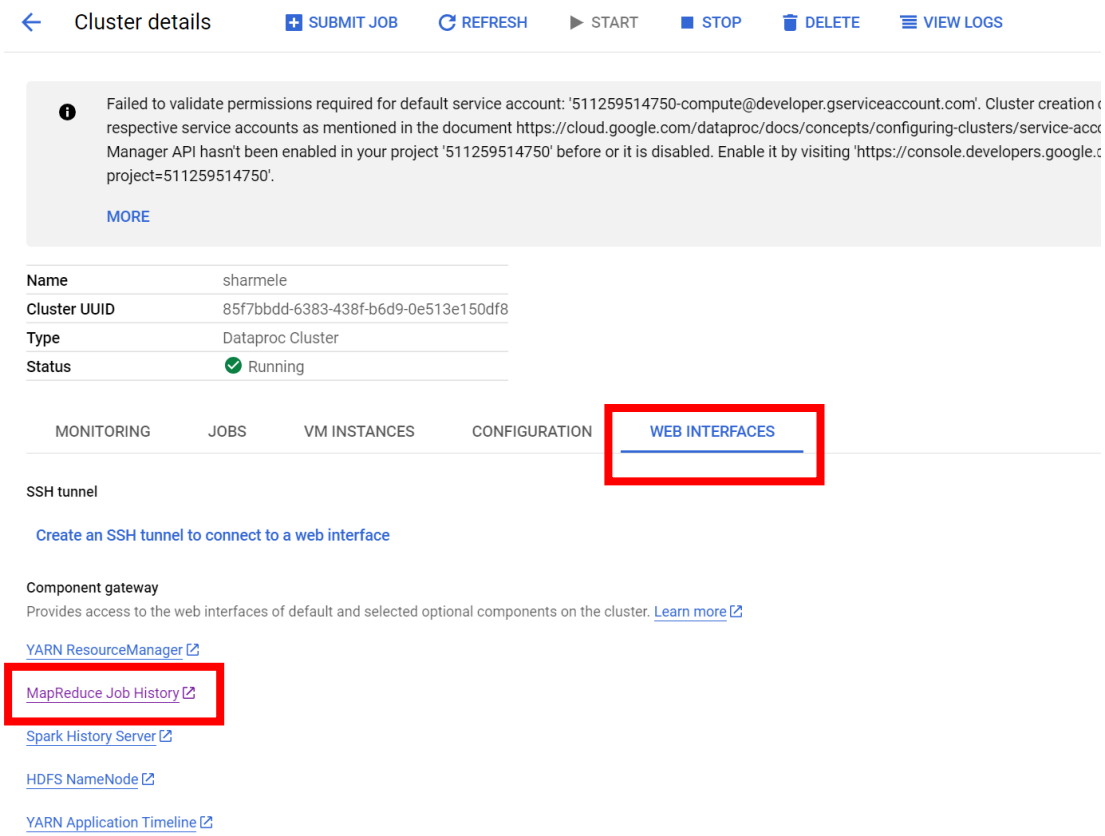

We could observe that there are 10 mappers and 1 reducer created for joining two tables and fetching student records with date of birth greater than '01/01/1995'.